

A/B testing posts and iterating

Week 5, Lecture 2 · Grow on X: From Zero to 10,000 Followers

Autumn 2026

What we'll cover

- Running the same idea through two different hooks to find the winner
- Sample sizes that work for solo creators (smaller than most assume)
- The weekly retrospective: keep, kill, or double down
- Avoiding p-hacking your own posting data

By the end of this lecture you can run a real A/B test on your next post, score the result, and feed the finding into next week's plan.

Same idea, two hooks

The concept of a post A/B test

- An A/B test for posts means publishing the same core idea twice, with one element changed
- The element you change is almost always the hook (first line)
- Everything downstream, the insight, the examples, the structure, stays the same
- The hook alone drives whether someone stops scrolling

Basilio (2023) on Dickie Bush: "He puts the idea out as a single tweet to validate resonance before committing resources."

What counts as one "element"

- Hook type: question vs. bold claim vs. statistic
- Framing: personal story opener vs. abstract principle opener
- Specificity: "I grew from 200 to 12k followers" vs. "most creators stall before 1k followers"
- Format: single tweet vs. first line of a thread

Change one element per test. If you change the hook and the format, you cannot tell which variable drove the result.

Writing two hooks for the same idea

A process that works in three steps:

1. Write the body of the post first. Lock it in.
2. Write five candidate hooks without filtering. Volume matters here.
3. Pick the two that are most different from each other in type or framing.

The most different pair gives you the most information. Two hooks that are subtly similar produce an inconclusive test.

Sample sizes for solo creators

Why "statistical significance" is the wrong frame

- Academic A/B tests need hundreds of conversions per variant to reach $p < 0.05$
- A solo creator posting twice a day has maybe 20-30 comparable posts per month
- That is not enough for textbook significance, and waiting for it means never learning
- The right frame is not significance testing, it is directional evidence plus replication

Chen (2018): "A rising curve tells you nothing predictive." A single win is also not predictive. Look for a pattern across five or more matched pairs.

A practical sample size for creators

- Run each hook variant at least three times before drawing a conclusion
- Three matched pairs (six posts total) gives you a directional read, not a proof
- A result that holds across all three pairs is strong enough to act on
- A result that splits 2-1 in favor of one variant means you need a fourth pair

The goal is a betting belief, not a published finding. Act on the belief, keep running the test, and update when new data arrives.

Controlling for confounders

Variables that can swamp your hook effect:

- **Time of day.** Post both variants at the same time on different days of the same week.
- **Day of week.** Weekday versus weekend can shift impressions by 30-50%.
- **News cycle.** A major event can make any post do better or worse.
- **Topic freshness.** If you already posted on this idea last week, response may be depressed.

Write down the posting conditions when you run the test. You cannot control for what you did not record.

The weekly retrospective

Keep, kill, or double down

Every Friday (or the last day you post each week), run a five-minute retrospective on the previous seven days of posts.

Classify each post into one of three buckets:

- **Keep:** performed at or above your baseline. Worth repeating the hook type and format.
- **Kill:** performed below baseline for the second or third time. Retire the hook pattern.
- **Double down:** performed well above baseline. Extract the pattern and replicate it next week.

What baseline means

- Your personal baseline is the median impressions-per-follower across all posts in the past 28 days
- Calculate it once per month and update it. Early in your account it will shift quickly.
- A post above 1.5x your median is worth studying. A post below 0.5x is worth retiring.
- Do not compare your numbers to other creators. Compare your numbers to your own past.

Rachitsky (2021): "Which metric, if it were to increase today, would most accelerate my business' flywheel?" The retrospective is where you find the answer.

The retrospective log

Keep a simple log. One row per post:

Date	Hook type	Format	Topic	Impr/follower	Verdict
Mon	Question	Thread	Sales	2.4x	Double
Tue	Statistic	Single	Sales	0.6x	Kill
Wed	Story	Thread	Ops	1.1x	Keep

After four weeks, the pattern is visible. After eight weeks, you have a content strategy grounded in your own data, not advice from someone else's account.

Avoiding p-hacking your own data

What p-hacking looks like for a solo creator

P-hacking means adjusting your analysis until the result tells the story you wanted. For a creator it sounds like:

- "That post flopped but the timing was bad, so I won't count it."
- "I'll look at engagements instead of impressions because this post got great replies."
- "My hook style is working, I just need to exclude the posts from that weird week."

Every exception you make narrows the dataset toward the conclusion you already believed.

The pre-commitment discipline

Before you post, write down two things:

1. The hypothesis: "I believe a question hook will outperform a bold claim hook for this topic."
2. The metric you will use to judge it: impressions per follower at 48 hours.

Commit to both before you see the results. This is the equivalent of pre-registering a study. It is the single most effective protection against motivated reasoning in self-analytics.

When to override your data

Your data can be wrong. Two legitimate reasons to override a test result:

1. **External event.** A major news story or platform outage clearly affected the posting window.
Note this in the log and mark the result as contaminated, not as a data point.
2. **Pattern conflict.** A result conflicts with a clear pattern across ten or more previous posts.
Treat it as an outlier, run the test again, and see if it replicates.

Everything else is motivated reasoning. If the result surprises you, that is the most valuable result you can get.

Take-aways

1. An effective post A/B test changes exactly one element, usually the hook. The body of the post stays constant.
2. Solo creators do not need statistical significance, they need directional evidence across at least three matched pairs.
3. The weekly retrospective classifies every post as keep, kill, or double down. After eight weeks this becomes a grounded content strategy.
4. Confirmation bias corrupts self-analytics. Pre-commit to your hypothesis and metric before you post.
5. Override your data only for external events or a clear pattern conflict. Surprise findings are the most useful ones.

"He puts the idea out as a single tweet to validate resonance before committing resources."

Jake Basilio on Dickie Bush, *Growth in Reverse* (2023)

What's next

- **Section (this week):** Pull 30 days of X analytics, classify your top 10 posts, double down on 2 patterns, retire 1.
- **Reading:** Week 5 reading, Analytics, experiments, and iteration at [/c/grow-on-x-26au/readings/wk05](#)
- **HW4 (30-day growth experiment):** Now out. Your experiment design is due next week.
- **Week 6:** Monetisation foundations, the creator stack that converts followers into revenue.