

Behavioral interviews are pattern matching

Week 7, Lecture 1 · Landing the Offer: SWE Job Search for Senior CS Students

Autumn 2026

What we will cover

- Why behavioral interviews exist and what interviewers are actually scoring
- Leadership principles as a public scoring rubric
- The STAR structure: situation, task, action, result
- The hiring bar and what "bar raiser" means in practice
- Selecting the right stories for the right signals

By the end of this lecture you can identify which leadership principle a question targets and map a story from your past to the STAR structure in under two minutes.

Part 1: why behavioral interviews exist

The behavioral interview is not a vibe check

- Every interviewer is assigned 2-3 specific leadership principles to evaluate
- They are looking for pattern-matched evidence of past behavior, not personality
- Spolsky (2006): the core question is "smart and gets things done" -- behavioral rounds test the second half
- Past behavior is the best predictor of future behavior: that is the theoretical basis

The interviewer has a scorecard before you walk in.

What the scorecard actually looks like

- Each interviewer takes one round with one or two assigned principles
- After the interview they write a detailed feedback packet: what was asked, verbatim quotes from the candidate, a hire/no-hire judgment
- The debrief reads every packet before a hiring decision is made
- Vague stories ("we decided as a team") score poorly because they give the committee nothing to evaluate

Orosz (2020): "Most interview prep materials focus almost exclusively on coding while ignoring the surrounding loop."

Part 2: leadership principles as a public rubric

Amazon's 16 principles are a teaching document

- Amazon publishes all 16 on amazon.jobs -- there is no hidden rubric
- Each principle has a name, a definition, and an implicit set of behaviors that constitute evidence
- Principles like Ownership, Bias for Action, and Deliver Results are the most commonly assessed for new-grad roles
- Other companies (Google, Meta, Microsoft) use different names but test the same underlying behaviors

Amazon Leadership Principles (2024): "We use our Leadership Principles every day, whether we're discussing ideas for new projects or deciding on the best way to solve a problem."

The five principles most tested at new-grad level

Principle	What it tests
Ownership	Did you treat problems as yours to solve?
Bias for Action	Did you act with incomplete information?
Deliver Results	Did you hit the goal despite obstacles?
Learn and Be Curious	Did you change your mind when wrong?
Earn Trust	Did you repair a damaged relationship?

Know the definition of each principle in your own words before your first interview.

Part 3: STAR structure

STAR: the four-part format every answer needs

- **Situation:** one or two sentences on context -- company, project, semester, team size
- **Task:** your specific responsibility or the decision that was yours to make
- **Action:** the three to five concrete steps you took; this is where most of the time goes
- **Result:** what happened, in numbers if possible; what you learned

Total time: under three minutes. If you run longer, you are over-explaining the Situation.

Where candidates lose points

- Too much Situation, not enough Action: the story becomes a movie trailer
- Plural subjects in Action ("we decided to...") hide individual ownership
- Result is vague: "it went well" or "the team was happy" is not a result
- The story does not connect back to the principle asked: the interviewer asked about Ownership but the story is about communication

Orosz (2023): "Practiced answers sound practiced -- and that is disqualifying."

Bad version vs. good version

"We had a production bug and the team worked on it. I helped debug and eventually we fixed it. The users were happy."

Strong answer (Ownership)

Weak answer (Ownership)

No individual action. No measurable result.
"We" throughout.

"A 3 AM alert woke me on a Saturday -- I was the only one with deploy access. I rolled back the build myself, filed the incident report, and ran the post-mortem the following Monday. Downtime went from two hours to 18 minutes compared to the prior incident."

Part 4: the hiring bar and story selection

What "the bar" means

- Every company sets a level: a story that would impress a new-grad bar versus a staff-engineer bar
- For new-grad roles the bar is: did you take initiative, work through ambiguity, and deliver something that mattered to someone?
- You do not need to have led a team of 20 or saved a production system
- You do need to have made a decision, not just followed instructions

Spolsky (2006): "It is much, much better to reject a good candidate than to accept a bad candidate."

The bar raiser role

- Amazon and some other companies include a designated "bar raiser" -- an interviewer from outside the hiring team
- Their job is to reject candidates the team wants but who do not meet the company-wide bar
- The bar raiser does not care about headcount pressure; they only score the candidate
- This matters because a team eager to hire you may be overruled
- Your stories need to hold up to a stranger's scrutiny, not just to the hiring manager who interviewed you

Choosing the right story for the right signal

- Each STAR story should map cleanly to one principle; forcing a story to cover three signals weakens it
- Story selection rule: the Action section must show a decision that was yours alone
- Prefer stories where you can name a specific number in the Result (latency dropped by 40%, user retention went up 15 points)
- Have at least two stories per principle you know will be tested -- interviewers follow up when they like a story

Your story bank is the deliverable for HW 5.



The behavioral interview is not a test of your storytelling. It's a test of whether you've actually been in hard situations and learned from them. Practiced answers sound practiced -- and that is disqualifying.

Take-aways

1. Behavioral interviews are structured: each interviewer scores one or two named principles, then writes a detailed packet.
2. Leadership principles are public documents. Reading them tells you exactly what the interviewer is looking for before you walk in.
3. The STAR structure is a contract: Situation sets context, Task names your responsibility, Action shows your decisions, Result demonstrates impact.
4. Individual ownership language in the Action section ("I decided," "I shipped," "I ran") is what separates scoreable stories from background noise.
5. Build a story bank now. Two stories per principle beats one polished story per principle.

What's next

- **Lecture 2 (this week):** The full interview loop -- phone screen, virtual on-site, debrief, what the committee reads
- **Section (this week):** STAR story bank workshop -- write 12 stories, peer-review three under timed conditions
- **Reading (this week):** Week 7 reading covers STAR in depth with a worked example and a mapping of 12 story sources
- **HW 3:** Due this week (recorded mock coding interview)
- **HW 5:** Out this week -- STAR story bank (12 stories)
- **Capstone:** Out this week -- application sprint, full mock loop, one unconventional move