

The honest dashboard

Week 1, Lecture 1 · LangoBee to \$10K MRR

Summer 2026 · Operating a Solo Language App from Zero Revenue

What we'll cover

- Why vanity metrics are not just misleading but actively harmful
- The three-A test: actionable, accessible, auditable
- De-alting: stripping founder, alt, and bot accounts from your counts
- Single source of truth: one queryable store beats many dashboards
- The intel repo: recording decisions so the next session can grep them

Vanity vs actionable metrics

The core problem

A vanity metric trends up regardless of whether the product is working.

- Total registered users: only goes up, includes bots, alts, and one-time visitors
- Total pageviews: rises whenever anyone opens a tab
- "103 profiles": the number on LangoBee's dashboard right now

The problem is not that these numbers lie. It is that teams believe them:

When numbers rise, teams credit themselves. When numbers fall, teams blame others.

The result is two private realities in the same founder's head, one optimistic, one pessimistic, neither grounded.

(Ries, 2009: startuplessonslearned.com)

The three-A test (Ries, 2009)

A metric is actionable if it passes three tests:

1. **Actionable:** a specific action caused the change. If you can't name the action, it's vanity.
2. **Accessible:** any team member can understand it without a data-science degree.
3. **Auditable:** it can be re-derived from raw data at any time.

Applying the test to LangoBee's dashboard:

Metric	Actionable	Accessible	Auditable	Verdict
103 total profiles	No	Yes	Yes	Vanity
4 de-alted WAU	Yes	Yes	Yes	Actionable
59% founder events	Yes	Yes	Yes	Actionable

What "actionable" looks like

Vanity metric: "103 registered profiles"

Actionable replacement: "4 de-activated weekly active users"

The replacement answers a different question:

- Not "how many people ever signed up" (only goes up)
- But "how many real humans used the product this week" (goes up when something works, down when something breaks)

The shift is always: **counts to rates, aggregates to per-user figures, totals to cohort slices.**

(Ries, 2009; Chen, 2015)

De-aling: making the numbers real

What de-alting means

De-alting removes three classes of accounts from your metrics:

1. **Founder accounts:** your own dogfooding, which is real usage but not customer behavior. At LangoBee, 59% of all events in the last 30 days come from one founder account.
2. **Alt accounts:** secondary accounts created for testing. They inflate activation rates (you always complete onboarding on a test account).
3. **Bot and spam registrations:** accounts that signed up programmatically. The June 1 week added 52 signups of which approximately 50 were bots.

A metric computed over all accounts measures activity in your org. A de-alted metric measures whether customers are getting value.

The June 2026 bot-signup spike

June 1 week: 52 signups. 1 activated.

Three signals identify the bots:

- **Email domains:** junk TLDs like a7gi.ru and host315.net that do not correspond to real email providers
- **Auth method:** password-only signups with no OAuth, which bots prefer because OAuth requires a real browser session
- **Post-signup events:** zero reading_events ever, not even a pageview on a text

A single compound filter removes them: exclude emails matching a known-junk-domain list OR accounts with zero reading_events within 7 days of signup.

Without this filter, the June cohort looks like 5x your normal acquisition week. With it, it looks like noise.

The de-alting filter in practice

```
-- Exclude from all metrics
WHERE email NOT IN ('hughgramelpacher@gmail.com')
      AND email NOT SIMILAR TO
          '%(a7gi\.ru|host315\.net|mailnull\.com)%'
```

In PostHog HogQL, the same filter runs on person properties:

```
AND person.properties.$email
    NOT IN ('hughgramelpacher@gmail.com')
AND NOT match(
    person.properties.$email,
    '(?i)(a7gi\\.ru|host315\\.net)'
)
```

Save this as a reusable person filter in PostHog. Every insight on your Ground truth dashboard applies it.

Single source of truth

Why multiple dashboards fail

Most founders have metrics scattered across:

- PostHog (events, funnels, cohorts)
- Supabase (user table, reading events, raw SQL)
- Vercel Analytics (traffic, but no API)
- Stripe (revenue)
- Memory ("I think it was around 4 active users")

When each tool shows a different number for "active users," you spend cognitive load reconciling rather than deciding. Worse, you unconsciously reach for whichever number supports your current belief.

One queryable store, one definition per metric, one dashboard.

The Ground truth dashboard

One PostHog dashboard named "Ground truth" with four pinned insights:

1. De-alted WAU (rolling 7 days)
2. De-alted MAU (rolling 30 days)
3. De-alted weekly signup trend (last 8 weeks)
4. Founder event share (sanity check, should stay below 60%)

The Supabase cohort table is the companion document: cohort-week rows with signup, activation, and return counts. Update it weekly by re-running the queries from section.

No other dashboard is authoritative. When a number conflicts, the Ground truth dashboard wins.

The intel repo

Decisions made without a record are repeated decisions. Each week, after reviewing the dashboard:

- Write one sentence: what changed, what caused it, what you will do next
- Store it in a dated file in the repo (e.g., `docs/intel/2026-06-09.md`)
- The file should be greppable: include the metric name, the value, and the decision

Example entry:

```
2026-06-09 | de-altd WAU: 4 | bot spike in Jun-1 cohort confirmed  
(52 signups, 1 real). Filter applied. No acquisition spend until  
WAU reaches 10 for two consecutive weeks.
```

A future session that asks "why did we not run ads in June?" can grep and find the answer.

When numbers rise, teams credit themselves. When they fall, teams blame others. Tracking vanity metrics produces divergent private realities that block any shared understanding of what is working.
(Paraphrased)

Eric Ries (2009)

Why Vanity Metrics Are Dangerous, Startup Lessons Learned

Take-aways

1. A vanity metric trends up regardless of product quality; the damage is not in being misled but in producing divergent private realities that block shared diagnosis.
2. The three-A test (actionable, accessible, auditable) is the filter: if you cannot name the action that caused a change, the metric is vanity.
3. De-alting removes founder events, alt accounts, and bot signups before any metric is computed. At LangoBee, not de-alting inflates raw WAU by approximately 15x.
4. One Ground truth dashboard with one definition per metric eliminates the reconciliation cost of scattered tools.
5. The intel repo turns decisions into greppable artifacts, so each session starts from truth rather than reconstructing it.

What's next

- **Reading:** "Ground truth and the de-alted dashboard" at </c/langobee-10k-mrr-26su/readings/wk01>
- **Lecture 2 (this week):** LangoBee's funnel, audited: AARRR mapped to real events, the full June 2026 snapshot, choosing a north-star metric
- **Section this week:** Rebuild the de-alted dashboard from raw HogQL and SQL
- **HW1 (going out this week):** The honest dashboard: publish your de-alted WAU, MAU, and cohort table