Quality is binary

Week 3, Lecture 1 · LangoBee to $10K MRR

Summer 2026 · Operating a Solo Language App from Zero Revenue

What we'll cover

- The broken-vacuum problem: what the Discord recruits were actually handed
- Must-work checklists: the minimum bar before inviting anyone
- Dogfooding vs the user's shoes: why founder eyes miss breakage
- Reading rageclicks: where session recordings tell the truth
- LangoBee's signals: 223 rageclicks, 34-second median session

The broken-vacuum problem

What the first users were handed

LangoBee's first real users came from Discord recruiting in May and June 2026.

The June 2026 snapshot:

- 4 real weekly active users
- 34-second median session
- 223 rageclicks in 30 days
- 1 share from 101 short-form completions

Word of mouth requires a product that works. If you hand someone a broken vacuum and ask them to tell a friend, they won't.

The broken-vacuum test

The broken-vacuum standard is not "are users delighted." It is:

> Does the product perform its core function, every time, for a new user, without instructions from the founder?

For LangoBee, the core function is: watch a short video, tap a word, see a definition.

If any of those three steps fails on an iPhone in Safari, the product fails the broken-vacuum test. Everything else is secondary.

(Seibel, 2019)

Why founders miss breakage

Three reasons the founder does not see broken-vacuum failures:

1. **Familiarity.** You know the workaround. The new user does not.
2. **Equipment.** Your test device is not the user's device. Safari on iPhone is not Chrome on MacBook.
3. **Context.** You sign in with saved credentials. The user starts from a Google result on a phone they have never used.

Dogfooding catches product-quality regressions. It does not catch onboarding breakage. Those are different problems.

Must-work vs nice-to-work

A must-work checklist has two columns only: pass and fail.

Must-work items for LangoBee:

- Feed loads (at least 5 videos, within 3 seconds on Wi-Fi)
- Captions render and are legible
- Word-tap popup opens (within 500 ms)
- Popup shows a definition, not a blank card
- Signup completes on iPhone Safari PWA and desktop Chrome
- Content does not run out within one session

If any item fails, the checklist fails. There is no partial credit.

(Seibel, 2019; Vohra, 2018)

Reading rageclicks

What a rageclick signals

A rageclick is three or more rapid clicks on the same element in quick succession.

It almost always means one thing: the user expected something to happen, and it did not.

LangoBee's June 2026 data: 223 rageclicks in 30 days across 31 monthly active users. That is roughly 7 frustrated interactions per active user per month.

The rageclick is not a bug report. It is a user who tried to tell you something was wrong and gave up before writing it down.

How to read session recordings

Five recordings will tell you more than 500 survey responses.

Protocol for each recording:

1. Note the URL when the rageclick fires
2. Note the element being clicked
3. Note how long the session lasts after the rageclick
4. Note whether the user reached word_lookup_opened at any point

If most sessions end within 10 seconds of the rageclick, the rageclick is likely the reason the session ended. If the session continues, the user pushed through.

(Nielsen, 2000)

Where the 34-second session dies

A 34-second median session means most users leave before they could possibly reach the aha moment.

Possible explanations, in order of diagnostic priority:

1. A broken UI element stops the user before the first video plays
2. Captions do not appear and the user has nothing to interact with
3. The word-tap affordance is not visible (no prompt, no highlight)
4. The user reached the word-tap popup but closed it and had nothing to do next

PostHog session recordings sorted by duration (20-50 seconds) will distinguish these. Watch three. Write one sentence per session.

> Launch something imperfect quickly, talk to users before and after, and resist falling in love with your MVP because it will need to change. (Paraphrased)

Michael Seibel, Y Combinator (2019)

How to Plan an MVP

The checklist is not a one-time event

Run the must-work checklist:

- Before inviting any new user (launch gate)
- After every deploy that touches the feed, reader, or auth flow
- After any dependency upgrade (Next.js, Supabase client, Capacitor)

Assign it to an agent. The checklist takes 20 minutes to run manually. An agent can run items F1-F4, W1-W3, C1-C2 against a headless browser in CI.

The items that require a real iPhone (S, P rows) run weekly, not per-deploy.

Take-aways

1. The broken-vacuum standard is binary: the product either performs its core function for a new user or it does not. There is no partial credit.
2. Founder dogfooding finds regressions. It does not find onboarding breakage. These are different problems requiring different protocols.
3. 223 rageclicks across 31 monthly active users is a signal that the product is actively breaking expectations for a majority of real users.
4. Session recordings sorted by duration reveal the specific moment where the 34-second session ends. Five recordings is enough.
5. The must-work checklist runs before every new-user invitation and after every deploy that touches core surfaces.

What's next

- **Reading:** "Week 3: The broken-vacuum bar and the guided first session" at `/c/langobee-10k-mrr-26su/readings/wk03`
- **Lecture 2 (this week):** Onboarding to first aha: time-to-value, the guided first session, activation defined
- **Section this week:** Run the must-work checklist on a real phone and laptop; file every defect as a GitHub issue
- **HW1 (due this week):** The checklist run is the submission; post results to the course Discord