

The weekly machine

Week 12, Lecture 1 · LangoBee to \$10K MRR

Summer 2026 · Operating a Solo Language App from Zero Revenue

What we'll cover

- The Monday report cadence: grade last week before planning this week
- Input metrics vs output metrics: the Amazon WBR architecture
- The langobee-intel repo as a CEO dashboard refreshed by agents
- Four agent-audited loops that close the operating system
- Why Sahil Lavingia's Gumroad story is a case study in absent cadence

Grade last week first

The sequencing discipline

Every Monday starts with one question: did last week's bets produce what you said they would?

Most founders answer this in reverse order. They open the build tool, then the calendar, and grade last week at the end of the day. If they grade it at all.

The Amazon Weekly Business Review inverts this:

- Metric review first, before any agenda item
- Grading the prior week is the agenda
- Planning the next week flows from the grade

(Chin, 2023)

What the Monday report looks like at LangoBee

The langobee-intel repo already has a weekly report structure. This week it upgrades from a narrative log to a lean WBR with input metrics.

The format, in order:

1. **North-star table:** five rows, filled before anything else
2. **Graded bets list:** green / yellow / red for each bet placed last Monday
3. **Theme-blocked calendar:** this week's committed build blocks
4. **Next bet placed:** one sentence, one metric threshold

Nothing in the build block opens until the north-star table is committed.

Why "grade last week first" is load-bearing

Skipping the grade produces two failure modes:

- **Motivated forgetting.** A bet that missed quietly disappears from the record. Three weeks later, you are still "testing" a channel with no data.
- **Sunk-cost persisting.** Because no explicit kill criterion was triggered, the bet continues past the point where evidence has falsified it.

Sahil Lavingia describes both at Gumroad: "The alarm bells were going off, but I was surrounded by people who believed in the mission." The grade did not happen because the cadence did not exist. (Lavingia, 2019)

Input metrics vs output metrics

The Amazon WBR distinction

Colin Bryar (former Bezos shadow) describes Amazon's WBR as built around two tiers of metrics:

- **Output metrics:** you observe them but cannot act on them directly. Revenue, MRR, churn rate. Useful for diagnosis.
- **Input metrics:** you can act on them this week. New push opt-ins, D1 return events, trial-start events. Useful for decisions.

The WBR discussion focuses entirely on inputs. If an output is off-track, the meeting asks: which input drove it, and what are you doing about that input this week?

(Chin, 2023)

LangoBee's input/output split

Output metric	Corresponding input metric
D1 return rate	New push opt-ins this week
Trial starts	Concierge onboarding sessions run
Real WAU	New activated users from the current channel
MRR	Trial-to-paid conversions this week
Shorts completion rate	New shorts published (fresh catalog)

An output that moves without a matching input change is a mystery. A mystery is not a strategy.

Why this matters for a solo founder

A venture-backed team has a growth pod, a PM, and a data analyst. Each owns an input metric.

A solo founder runs all of them. The risk is not insufficient data: it is too much data with no structure for action.

The input metric discipline answers: what is the one thing you will do this week that, if it moves, explains the output you want?

Pick one. Act on it. Grade it Monday.

The intel repo as CEO dashboard

What the langobee-intel repo already is

The weekly report in langobee-intel is already a CEO dashboard in structure:

- Version-controlled: every week's grade is queryable by a future session
- Agent-refreshable: PostHog queries and Supabase counts can be run by an agent and committed automatically
- Greppable: a cold session can find the D1 trend, the channel history, and the kill criteria without reading 12 weeks of narrative

The upgrade this week is not a new repo. It is making the north-star table the first thing in every entry.

The four agent-audited loops

A loop is closed when an agent can verify it without human intervention. Four loops should be agent-auditable by end of week 12:

1. **Metric pull:** PostHog WAU + D1 cohort + Supabase trial count, committed to intel every Monday
2. **Bot filter:** de-aleted query runs automatically; the intel report never includes founder or bot events
3. **Kill-criteria check:** the bets list is compared against the kill thresholds written at the start of the experiment
4. **Bet placement:** this week's bet is written to intel before the build block opens on Monday

None of these requires a human to do the math. They require a human to read the output and make a decision.

The WBR discussion stays on input metrics, the things you can control, not the output metrics you can only observe. (Paraphrased)

Colin Bryar, via Cedric Chin (2023)

Colin Bryar on the practice of Amazon's Weekly Business Review

What the Gumroad case teaches

Sahil Lavingia's 2019 essay reconstructs three decision points where a working cadence would have changed the outcome:

- A growth plateau that should have triggered a pivot-or-persevere review
- An acquisition offer evaluated without pre-committed criteria
- A fundraising round pursued after clear plateau signals

None of these required more data. They required a weekly meeting where the data was graded against pre-written conditions.

The cadence is not the cure. But its absence is reliably the cause. (Lavingia, 2019)

Take-aways

1. Grade last week before planning this week. The Monday cadence is load-bearing: without it, kill criteria go untriggered and sunk-cost bets persist.
2. Every output metric has a corresponding input metric you can act on this week. Run your WBR on inputs, not outputs.
3. The langobee-intel repo is already a CEO dashboard. Upgrading it to lean WBR format requires one structural change: north-star table first, committed before the build block opens.
4. Four loops should be agent-auditable: metric pull, bot filter, kill-criteria check, and bet placement.
5. The absence of a weekly cadence, not the presence of bad luck, is what Sahil Lavingia describes as the root cause at Gumroad.

What's next

- **Reading:** "Week 12: The \$10K MRR operating plan" at </c/langobee-10k-mrr-26su/readings/wk12>
- **Lecture 2 (this week):** From here to \$10K: sequenced bets, channel math, and the 12-month view
- **Section this week:** Capstone working session: present the operating plan and red-team it against the kill criteria
- **Capstone due this week:** The \$10K MRR operating plan with sequenced bets, channel math, and kill criteria