From here to $10K

Week 12, Lecture 2 · LangoBee to $10K MRR

Summer 2026 · Operating a Solo Language App from Zero Revenue

What we'll cover

- The sequencing argument: why retention comes before the first dollar, and the dollar comes before the channel
- Channel math: the arithmetic that constrains every growth plan
- The Max MRR formula as a leading indicator
- Five self-deceptions to audit before writing the operating plan
- The 12-month view with explicit kill criteria

The sequencing argument

Three things in order, not in parallel

The most common mistake at the $0-to-$10K stage is running all three bets in parallel:

- Retention work (improve D1 return rate)
- Revenue work (get the first trial started)
- Channel work (find a repeatable acquisition path)

A solo founder does not have the bandwidth to run three bets in parallel without all three running at half-effort. The correct sequence is one at a time, with a clear evidence threshold for moving to the next.

(Walling, 2015)

Bet 1: the retention bar

LangoBee's real cohorts: May 18 week (15 signups, 12 activated, 6 returned later), May 25 week (19 signups, 7 activated, 3 returned). Overall: approximately 24% ever-return on de-alted real users.

The retention bar for Bet 1: D1 return rate of 30% or above on a real cohort of at least 20 users, sustained for two consecutive weeks.

Below this bar, acquisition dollars are wasted. A user who does not come back on day 1 does not start a trial.

The course's retention phase (weeks 3-5) built the measurement system for this bar. Bet 1 is not closed until the data confirms it.

(June 2026 snapshot; Balfour, 2013)

Bet 2: the first dollar

Zero trial_started events ever. Stripe checkout, trial mechanics, and the webhook are all wired. No human has entered the funnel.

Bet 2 is not an engineering problem. It is a sales problem: manually recruiting a user who is willing to start a trial.

The evidence threshold: at least one trial started, converted to paid, and held for 30 days. Not a lifetime deal. Not a gifted subscription. A user who went through the paywall independently.

This is the Stair Step Method's first stair: one paying customer, acquired deliberately. (Walling, 2015)

Bet 3: one repeatable channel

Acquisition attribution (June 2026): direct 174 people (Discord manual recruiting), google.com 21, t.co 5.

"Direct" is not a channel. It is the founder doing sales by hand. That is correct for Bet 2. It is not a channel for Bet 3.

A channel is repeatable: you can go back and do it again with a predictable result. Three independent repetitions with consistent cost-per-signup and signup-to-activation rates, run without the founder's personal attention, is the evidence threshold for Bet 3.

(June 2026 snapshot)

Channel math

Building backward from $10K MRR

$10K MRR at $7.99/mo requires approximately 1,252 active paying subscribers.

Work backward:

```
$10,000 MRR / $7.99 ARPA = 1,252 subscribers
At 5% monthly churn: need ~63 new subscribers/month just to hold flat
At 25% trial-to-paid: need 252 trial starts/month
At 3% signup-to-trial: need 8,400 signups/month
```

These are not targets. They are constraints. If any conversion rate in the chain is unmeasured, the plan is fiction.

(Cohen, 2013; Walling, 2015)

The Max MRR formula as a ceiling check

Jason Cohen derives the formula: Max MRR equals new MRR divided by the cancellation rate.

If new MRR per month is $200 and the monthly cancellation rate is 5%, Max MRR is $4,000: a hard ceiling, visible months before the plateau materializes.

At LangoBee's current new MRR of $0 and any non-zero churn, Max MRR is $0 regardless of acquisition spend. The formula makes this explicit: the ceiling is a function of both acquisition and retention, not acquisition alone.

(Cohen, 2025)

Where the higher return is

- Max MRR at $200 new MRR: $10,000
- Max MRR: $8,000

(Cohen, 2025; Rachitsky, 2022)

Improving churn from 5% to 2%:

Doubling acquisition to $400 new MRR at 5% churn:

Retention improvement produces more MRR gain than doubling acquisition spend at this stage. This is the arithmetic case for doing Bet 1 before Bet 3.

The five self-deceptions

Self-deception 1: vanity count

The June 1 week showed 52 apparent signups. One was real.

Any operating plan that uses "total registered users" as a growth metric is off by approximately 50x at the current bot-to-human ratio. The correct input is real activated users on de-alted, domain-filtered cohorts.

Check: restate every user count in your plan as a cohort-based rate. If the rate version does not support the plan, the plan needs to change.

(June 2026 snapshot; Ries, 2009)

Self-deception 2: bot-signup inflation

Bot signups have junk-domain tells: a7gi.ru, host315.net. They exhibit a signature: password signups, exactly 1 activation event, never returning.

An operating plan built on raw Supabase user counts includes this population. The de-alted filter built in week 1 must be applied to every metric in the plan, not just the ones labeled "retention."

Self-deception 3: founder dogfooding inflation

59% of all PostHog events in the June 2026 snapshot were generated by the founder account. A median session of 34 seconds across real users looks longer when blended with founder sessions.

Any engagement figure that includes founder activity is inflated. The correct test: pull every engagement metric with the founder account, alt accounts, and team accounts excluded. If the number drops by more than 20%, the unfiltered number was fiction.

(June 2026 snapshot)

Self-deceptions 4 and 5: spike and channel confusion; channel math avoidance

Self-deception 4: one spike is a channel. LangoBee's one Show HN post, one creator mention, or one viral Reddit thread each produced a spike. A spike is a single event with unrepeatable conditions. It is not evidence that the channel is viable. Three independent repetitions with consistent conversion rates is the minimum for "channel."

Self-deception 5: not doing the channel math. The backward-from-$10K arithmetic in the previous section is uncomfortable because it makes the required traffic volume concrete. Founders who avoid this math are not avoiding discouragement; they are avoiding the only tool that distinguishes a viable plan from a hope.

> Kill criteria are written before the bet starts, as a state plus a date, so the decision to quit is made by the clear-eyed version of you rather than the one in the middle of losing. (Paraphrased)

Annie Duke (2023)

Author Talks: What poker pro Annie Duke can teach you about quitting on time (McKinsey)

The 12-month view

Three bets, in sequence, with explicit kill criteria:

Bet	Evidence threshold	Kill criterion	Fallback
Retention	D1 rate 30%+ for 2 cohorts	Below 20% after 8 weeks of hooks	Ship concierge onboarding, re-measure
First dollar	1 trial started and converted	0 trials from 50 real signups by week 16	Lower paywall, test free-trial variant
Channel	3 repetitions, consistent CPR	CPR above $50 after 3 runs	Test next Bullseye channel

Kill criteria are written before the bet is placed, not after the data comes in.

(Duke, 2023; Walling, 2015)

This week's course grade

The capstone asks you to grade twelve weeks against the course's mission:

- **Foundations (wk 1-2):** what did the thesis discipline prove, kill, and leave open?
- **Retention (wk 3-5):** what did the D1 work prove, kill, and leave open?
- **Monetization (wk 6-8):** what did the MRR model prove, kill, and leave open?
- **Distribution (wk 9-11):** what did the channel work prove, kill, and leave open?
- **Synthesis (wk 12):** what does the operating plan commit to as the next bet?

The grade is not a reflection essay. It is an evidence audit: name the specific number that supports each claim.

Take-aways

1. The correct sequence is retention bar first, first dollar second, one repeatable channel third. Running all three in parallel on solo-founder bandwidth means all three run at half-effort.
2. The Max MRR formula (new MRR divided by cancellation rate) shows the ceiling months before the plateau. At $0 new MRR, the ceiling is $0 regardless of churn.
3. The five self-deceptions (vanity count, bot inflation, dogfooding inflation, spike-as-channel, channel-math avoidance) have all appeared in LangoBee's numbers. Audit the plan against each before the capstone.
4. Kill criteria are written before the bet is placed in state-plus-date format. A bet without a kill criterion is not a bet; it is an open loop.
5. The 12-month view is traceable to specific PostHog and Stripe observations, not aspirations.

What's next

- **Section this week:** Capstone working session: present the operating plan, red-team it against the kill criteria, run the Monday cadence once end to end
- **Capstone deliverable:** $10K MRR operating plan: three sequenced bets, channel math, kill criteria, and 12-month timeline
- **Course grade:** five phases, each with one proven claim, one killed hypothesis, one open question